



Using Sun Grid Engine and Globus to Schedule Jobs Across a Combination of Local and Remote Machines

By Geoff Cawood and Paul Graham
Edinburgh Parallel Computing Centre (EPCC)

Telephone: +44 131 650 5120

Email: geoffc@epcc.ed.ac.uk

► EPCC

- Edinburgh Parallel Computing Centre
- Part of the University of Edinburgh in Scotland
- “A technology-transfer centre for high-performance-computing”
- <http://www.epcc.ed.ac.uk>


► Project Personnel

- Terry Sloan
 - Project leader - tms@epcc.ed.ac.uk
- Thomas Seed
- Ali Anjomshoaa
- Geoff Cawood
- Paul Graham

- ▶ A collaboration between EPCC and Sun
 - Called simply ‘Sungrid’ within EPCC
 - 57 person months of EPCC effort
 - Significant staff contribution from Sun


► Project scenario

- Two collaborating sites A and B both have some machines
- Both sites run Sun Grid Engine (SGE) to schedule jobs
- Local demand for machines is variable
 - Sometimes it exceeds supply
 - Other times machines lie idle



► Ideal Situation

- If sites A and B could expose their machines to each other across the internet through SGE...
- Both sites could enjoy through-put efficiency improvements
- Large gains when one site is busy while the other is idle




► Final goal

- Develop a job-scheduler based on SGE which can schedule jobs across a combination of local and remote machines
- Globus will provide a secure means of running jobs on remote sites

- ▶ Self-education and code analysis
 - Current phase
 - Learning about SGE internals and Globus
 - Documenting relevant sections using UML
- ▶ Requirements
- ▶ Design
- ▶ Implementation
- ▶ Final test

- ▶ Meet new colleagues
- ▶ Learn more about the internals of SGE
- ▶ Find out about any other attempts to extend SGE functionality over multiple sites
- ▶ Solicit feedback on the solution approaches we might adopt
 - Three approaches outlined in the remaining slides



► Outline


- SGE+ is SGE enhanced to work with remote machines
- Both sites can run SGE+
- Change scheduler to take local/remote issues into account
- Modify daemons (execd, commd etc.) to work across internet?
- Or use globus calls to fork jobs and query machine and job status?

► Pros

- Simple approach (?)
- Usability – nothing new for users to learn

► Cons

- Low quality scheduling decisions (?)
 - Neither SGE+ system knows exactly what jobs are running on a machine
 - May be a time-lag in getting query results back from remote site
- Duplicated scheduling effort
 - Both SGE+ systems are trying to do the same thing - wasteful
- Poor scalability
 - With N sites, might need N execd daemons per machine
- Queue control – awkward to turn off remote access temporarily
 - Have to explicitly ask other sites to stop forking jobs
- Remote load management tool is by-passed
 - Remote administrator might not want lots of forked jobs



► Outline


- Similar to previous approach but one site (A) nominated as master
 - All jobs from all sites are submitted through the master
- Site A runs SGE+ as described in previous approach
- Users at site B call `globus_run` to submit jobs to the enhanced SGE+ running at site A

► Pros

- Good scheduling decisions due to accurate information
- Under the remote daemons scheme:
 - Site B users might still be able to enjoy SGE job monitoring
 - Site B administrator might still have convenient queue control

► Cons

- Need back-up server/site
 - Users at B will not want to wait for A to come back up
- Political issues
 - Site B administrator loses control – may not be acceptable
- Need very high degree of co-operation
 - Correct software applications installed and maintained
 - User-support
 - Accounting issues



► Outline

- Build a new higher-level scheduler that schedules across globus job-recipients rather than individual machines or queues

► Outline continued

- Define a ‘globus job-recipient’ as a machine or load-management tool that can:
 - Receive jobs via globus and execute them
 - Respond to queries (sent via globus) about machine and job status
- Modify SGE to provide such a job-recipient interface
 - If it doesn’t already... (qstat may be enough?)
- Introduce a new higher-level Hierarchical Scheduler (HS) to schedule across job-recipients
 - Each site has its own HS
 - Only users at site A can submit jobs to the HS at A
 - Only users at site B can submit jobs to the HS at B
 - When a job arrives at a HS, it queries the job-recipients then decides which one to send the job to

► Pros

- Problem decomposition
 - Divide and conquer
 - Usually a good idea
- Should give reasonable efficiency improvements
- Local scheduling decisions made locally, as today
- Queue access controlled locally
 - Administrator can turn queues on/off easily through SGE, as today
- Extensible
 - Just make each HS aware of new job-recipients
- No ‘behind-the-scenes’ job-forking
 - Jobs arrive at a site through the intended entry point so SGE has full knowledge about what jobs are running, as today

► Cons

- Large effort to build new HS tool and front-end
- Scheduling decision duplication
 - HS decision algorithm may be very similar to what SGE does already
- HS scheduling is coarser-grained
 - So decisions might not be as good as centralised scheduling
- HS user-interface uncertain
 - What information on job progress would the HS user see?
- Political issues
 - This may not be what our project proposers want/expect

- ▶ If you have any comments on these approaches please come and talk to us!
- ▶ We are very interested to hear your opinions
- ▶ Danke schön!